



KDD2016

22nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining
August 13 - 17, 2016 | San Francisco, California

Streaming Analytics **Tutorial**

Ashish Gupta (LinkedIn)
Neera Agarwal

Streaming Analytics

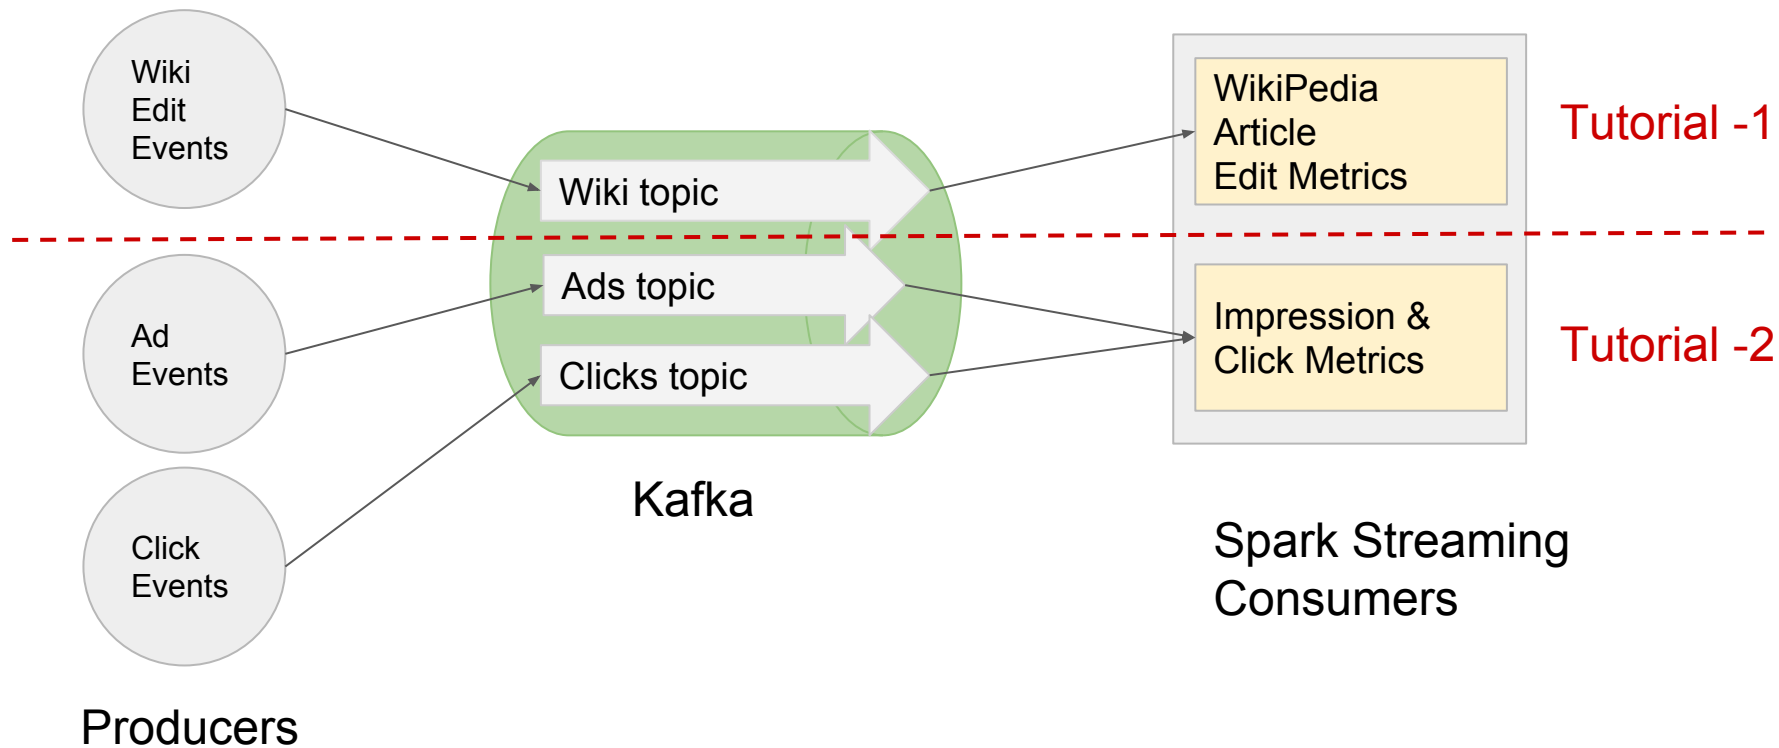
Before doing this tutorial please read the main presentation.

<http://www.slideshare.net/NeeraAgarwal2/streaming-analytics>

Technical Requirements

- OS: MAC OS X
- Programming Language: Scala 2.10.x
- Open source software used in tutorial: Kafka and Spark 1.6.2

Tutorials



Step 1: Installation

In the conference we provided USB stick with the environment ready for the tutorials. Here are the instructions to create your own environment:

1. Check Java: `java -version`

`java version "1.8.0_92"` (Note: Java 1.7+ should work.)

2. Check Maven: `mvn -v`

If not installed, check instructions in the **Additional slides** at the end.

Step 1: Installation

3. Install Scala

```
curl -O http://downloads.lightbend.com/scala/2.10.6/scala-2.10.6.tgz
```

```
tar -xzf scala-2.10.6.tgz
```

Set scala home to the path pointing to scala-2.10.6 folder. For example on Mac:

```
export SCALA_HOME=/Users/<username>/scala-2.10.6
```

```
export PATH=$PATH:$SCALA_HOME/bin
```

4. Install Spark

```
curl -O http://d3kbcqa49mib13.cloudfront.net/spark-1.6.2-bin-hadoop2.6.tgz
```

Step 1: Installation

5. Install Kafka

```
curl -O http://apache.claz.org/kafka/0.10.0.0/kafka\_2.10-0.10.0.0.tgz
```

```
tar -xzf kafka_2.10-0.10.0.0.tgz
```

6. Download tutorial

<https://github.com/NeeraAgarwal/kdd2016-streaming-tutorial>

Step 2: Start Kafka

Start a terminal window. Start Zookeeper

```
> cd kafka_2.10-0.10.0.0  
> bin/zookeeper-server-start.sh config/zookeeper.properties
```

Wait for zookeeper to start. Start another terminal window and start Kafka

```
> cd kafka_2.10-0.10.0.0  
> bin/kafka-server-start.sh config/server.properties
```


Tutorial 1:

Bot and Human edit counts on Wikipedia Edit stream

Step 1: Listening to Wikipedia Edit Stream

Start a new terminal window. Run Wikipedia Connector

```
> cd kdd2016-streaming-tutorial
```

```
> java -cp target/streamingtutorial-1.0.0-jar-with-dependencies.jar example.WikipediaConnector
```

After some messages, stop using CTRL-C. We will run it again after writing streaming code.

Does not run, build package and try again...

```
> mvn package
```

Step 1: Wikipedia Stream message structure

[[-acylglycerol O-acyltransferase]] MB

https://en.wikipedia.org/w/index.php?diff=733783045&oldid=721976415 * BU RoBOT * (-1) /* References
*/Sort into more specific stub template based on presence in [[Category:EC 2.3]] or subcategories (Task
25)

[[City Building]] <https://en.wikipedia.org/w/index.php?diff=733783047&oldid=732314994> * Hmains * (+9)
refine category structures

[[Wikipedia:Articles for deletion/Log/2016 August 10]] **B**

<https://en.wikipedia.org/w/index.php?diff=733783051&oldid=733783026> * Cyberbot

Fields: title, flags, diffUrl, user, byteDiff, summary

Flags (2nd field): 'M'=Minor, 'N' = New, '!' = Unpatrolled, 'B' = Bot Edit

WikiPedia Stream pattern = `"\\[[\\[(.*)\\]]\\]\\s(.*)\\s(.*)\\s\\|*\\s(.*)\\s\\|*\\s\\|\\|+?(\\.\\|d*)\\]\\]\\s(.*)".r`

Spark: RDD

An RDD is an immutable distributed collection of objects. Each RDD is split into multiple [partitions](#), which may be computed on different nodes of the cluster.

RDDs can contain any type of objects, including Python, Java, Scala or user-defined classes.

RDDs offer two types of operations:

- [Transformations](#) construct a new RDD from a previous one.
- [Actions](#) compute a result based on an RDD, and either return it to the driver program or save it to an external storage system.

Spark: DStream

DStream is a sequence of data arriving over time.



Internally, each DStream is represented as a sequence of RDDs arriving at each time step. DStreams offer two types of operations:

- **Transformations** yield a new DStream.
- **Output operations** write data to an external system.

Step 2: Write Code

In [kdd2016-streaming-tutorial](#)

Change code in `src/main/scala/example/WikipediaStreaming.scala` file. Use your favorite editor.

```
val lines = messages.foreachRDD { rdd =>
```

```
    // ADD CODE HERE
```

```
}
```

Note: In the conference participants were asked to write the code while in github repository full code is provided.

Step 2: Continued

rdd =>

```
val linesDF = rdd.map(row => row._2 match {
```

```
  case pattern(title, flags, diffUrl, user, byteDiff, summary) => WikiEdit(title, flags, diffUrl, user,  
byteDiff.toInt, summary)
```

```
  case _ => WikiEdit("title", "flags", "diffUrl", "user", 0, "summary")
```

```
}).filter(row => row.title != "title").toDF()
```

Step 2 Continued

// Number of records in 10 second window.

```
val totalCnt = linesDF.count()
```

// Number of bot edited records in 10 second window.

```
val botEditCnt = linesDF.filter("flags like '%B%'").count()
```

// Number of human edited records in 10 second window.

```
val humanEditCnt = linesDF.filter("flags not like '%B%'").count()
```

```
val botEditPct = if (totalCnt > 0) 100 * botEditCnt / totalCnt else 0
```

```
val humanEditPct = if (totalCnt > 0) 100 * humanEditCnt / totalCnt else 0
```


Step 3: Build Program

Start a new terminal window.

```
> cd kdd2016-streaming-tutorial
```

```
> mvn package
```

Step 4: Run Programs

Run `WikiPediaConnector` in terminal window. It starts receiving data from `WikiPedia IRC` channel and writes to `Kafka`.

```
> java -cp target/streamingtutorial-1.0.0-jar-with-dependencies.jar  
example.WikiPediaConnector
```

Run `WikiPediaStream` in a new terminal window.

```
> cd kdd2016-streaming-tutorial  
  
> ../spark-1.6.2-bin-hadoop2.6/bin/spark-submit --class  
example.WikiPediaStreaming  
target/streamingtutorial-1.0.0-jar-with-dependencies.jar
```

Output

BotCount	BotPct	HumanCount	HumanPct
8	36	14	63

BotCount	BotPct	HumanCount	HumanPct
6	21	22	78

BotCount	BotPct	HumanCount	HumanPct
5	21	18	78

BotCount	BotPct	HumanCount	HumanPct
8	29	19	70

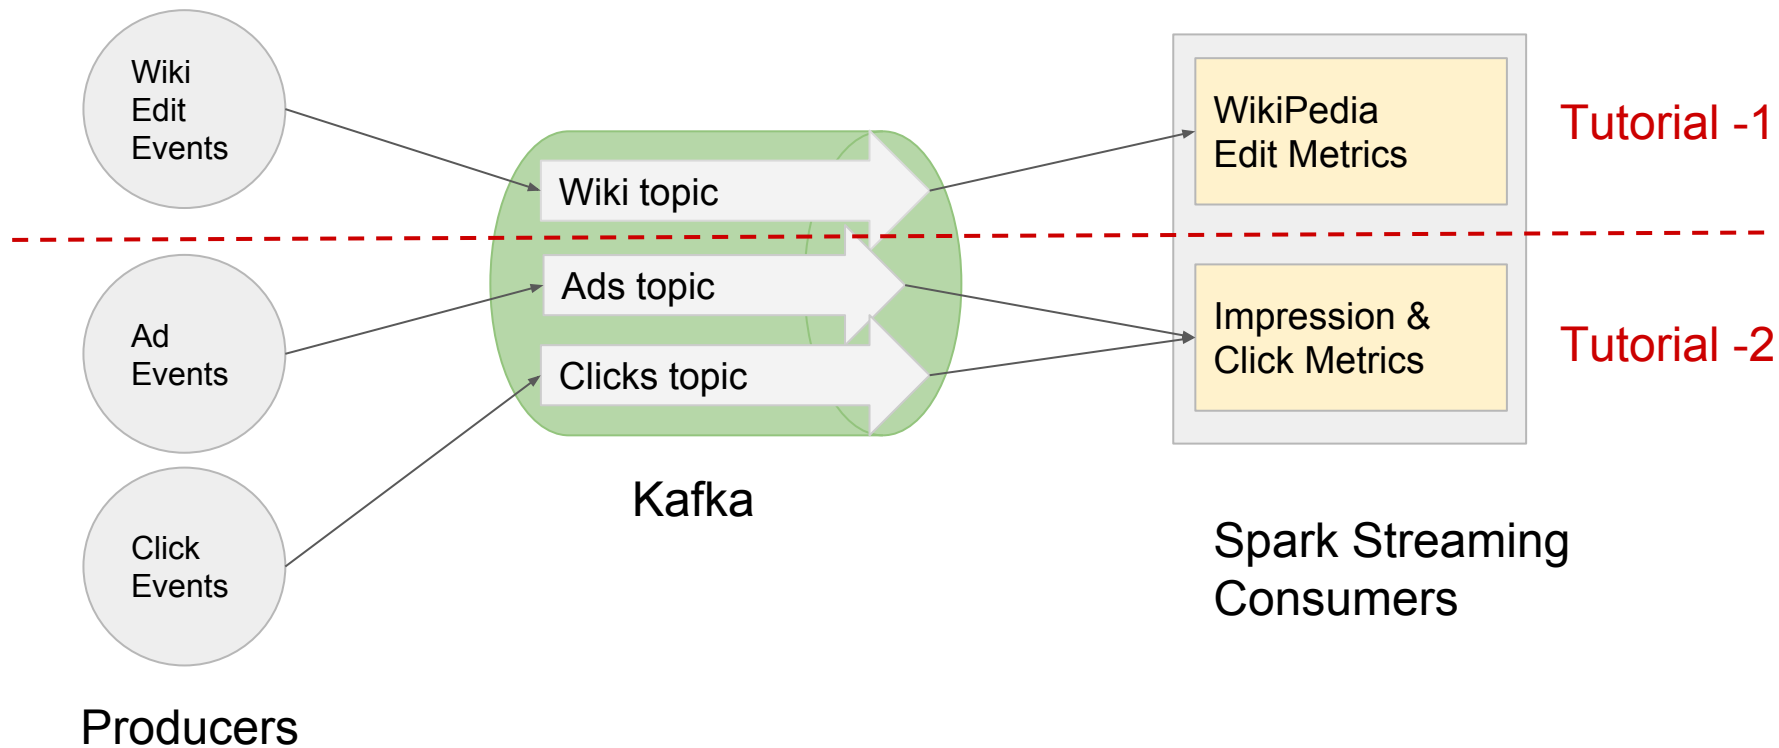
BotCount	BotPct	HumanCount	HumanPct
3	16	15	83

BotCount	BotPct	HumanCount	HumanPct
6	25	18	75

Tutorial 2:

Impression Click metrics on Ad and Click streams

Tutorials



Step 1: Listening to Ads and Clicks stream

Run program to replay Ad and Click Events from file

```
> cd kdd2016-streaming-tutorial
```

```
> java -cp target/streamingtutorial-1.0.0-jar-with-dependencies.jar example.AdClickEventReplay
```

After some messages, stop using CTRL-C. We will run it again after writing streaming code.

Does not run, build package and try again...

```
> mvn package
```

Step 1: Ad and Click Event message structure

Ad Event:

QueryID, AdId, TimeStamp: 6815, 48195, 1470632477761

Click Event:

QueryID, ClickId, TimeStamp: 6815, 93630, 1470632827088

Join on QueryId, show metrics by Ad Id.

Step 2: Write Code

In [kdd2016-streaming-tutorial](#)

Change code in [src/main/scala/example/AdEventJoiner.scala](#) file.

```
val adEventDStream = adStream.transform( rdd => {  
  
    rdd.map(line => line._2.split(";")).  
  
        map(row => (row(0).trim.toInt, AdEvent(row(0).trim.toInt, row(1).trim.toInt, row(2).trim.toLong)))  
  
    })  
  
// ADD CODE HERE..
```

Note: In the conference participants were asked to write the code while in github repository full code is provided.

Step 2 Continued

//Connects Spark Streaming to Kafka Topic and gets DStream of RDDs (click event message)

```
val clickStream = KafkaUtils.createDirectStream[String, String, StringDecoder, StringDecoder](ssc,  
kafkaParams, clickStreamTopicSet)
```

//Create a new DStream by extracting kafka message and converting it to DStream[queryId, ClickEvent]

```
val clickEventDStream = clickStream.transform{ rdd =>  
  
  rdd.map(line => line._2.split(",")).  
  
    map(row => (row(0).trim.toInt, ClickEvent(row(0).trim.toInt, row(1).trim.toInt, row(2).trim.toLong)))  
  
}
```

Step 2 Continued

// Join adEvent and clickEvent DStreams and output DStream[queryId, (adEvent, clickEvent)]

```
val joinByQueryId = adEventDStream.join(clickEventDStream)
```

```
joinByQueryId.print()
```

// Transform DStream to DStream[adId, count(adId)] for each RDD

```
val countByAdId = joinByQueryId.map(rdd => (rdd._2._1.adId, 1)).reduceByKey(_+_)
```

Step 2 Continued

// Update the state [adId, countCumulative(adId)] by values from the next RDDs

```
val updateFunc = (values: Seq[Int], state: Option[Int]) => {
```

```
    val currentCount = values.sum
```

```
    val previousCount = state.getOrElse(0)
```

```
    Some(currentCount + previousCount)
```

```
}
```

```
val countByAdIdCumm = countByAdId.updateStateByKey(updateFunc)
```

// Transform (key, value) pair to (adId, count(adId), countCumulative(adId))

```
val ad = countByAdId.join(countByAdIdCumm).map {case (adId, (count, cumCount)) => (adId, count, cumCount)}
```

Step 2 Continued

//Print report

```
ad.foreachRDD( ad => {  
  
    println("%5s %10s %12s".format("AdId", "AdCount", "AdCountCumm"))  
  
    ad.foreach( row => println("%5s %10s %12s".format(row._1, row._2, row._3)))  
  
})
```

Step 3: Build Program

```
> cd kdd2016-streaming-tutorial
```

```
> mvn package
```

Step 4: Run Programs

Run `AdClickEventReplay` in terminal window. It reads data from Ad and Click event files writes to Kafka.

```
> cd kdd2016-streaming-tutorial
```

```
> java -cp target/streamingtutorial-1.0.0-jar-with-dependencies.jar  
example.AdClickEventReplay
```

Run `AdEventJoiner` in a new terminal window.

```
> cd kdd2016-streaming-tutorial
```

```
> ../spark-1.6.2-bin-hadoop2.6/bin/spark-submit --class example.AdEventJoiner  
target/streamingtutorial-1.0.0-jar-with-dependencies.jar
```

Output

AdId	AdCount	AdCountCumm
4	4	35
0	9	35
1	4	34
7	3	23
6	5	27
3	2	30
8	7	27
9	4	29
5	3	26
2	6	27

Contact Us

Ashish Gupta - ahgupta@linkedin.com

<https://www.linkedin.com/in/guptash>

Neera Agarwal - neera8work@gmail.com

<https://www.linkedin.com/in/neera-agarwal-21b9473>

Additional Notes - Install Java (Mac 10.11)

- `java -version`

```
java version "1.8.0_92"
```

If java does not exist, try

In `.bash_profile` add `export JAVA_HOME=$(/usr/libexec/java_home)`

(Install Java - https://java.com/en/download/help/mac_install.xml)

Additional Notes - Install Maven

- Check Maven on a terminal window
 - `mvn -v`
 - Apache Maven 3.2.5+
- Install Maven
 - `brew install maven`

OR if you do not have brew then do:

1. `curl -O`
<http://mirror.nexcess.net/apache/maven/maven-3/3.3.9/binaries/apache-maven-3.3.9-bin.tar.gz>
2. `tar -xzf apache-maven-3.3.9-bin.tar.gz`
3. In `.bash_profile`: `export PATH=/<path>/apache-maven-3.3.9/bin:$PATH`