



# KDD2016

22nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining  
August 13 - 17, 2016 | San Francisco, California


## Streaming Analytics

Ashish Gupta (LinkedIn)

Neera Agarwal

# What is a Data Stream

- Unbounded Data
- Data arriving continuously at high rate
- Too large to first store and then process
- Need to be processed in one pass
- May display Temporal Locality - patterns may evolve over time



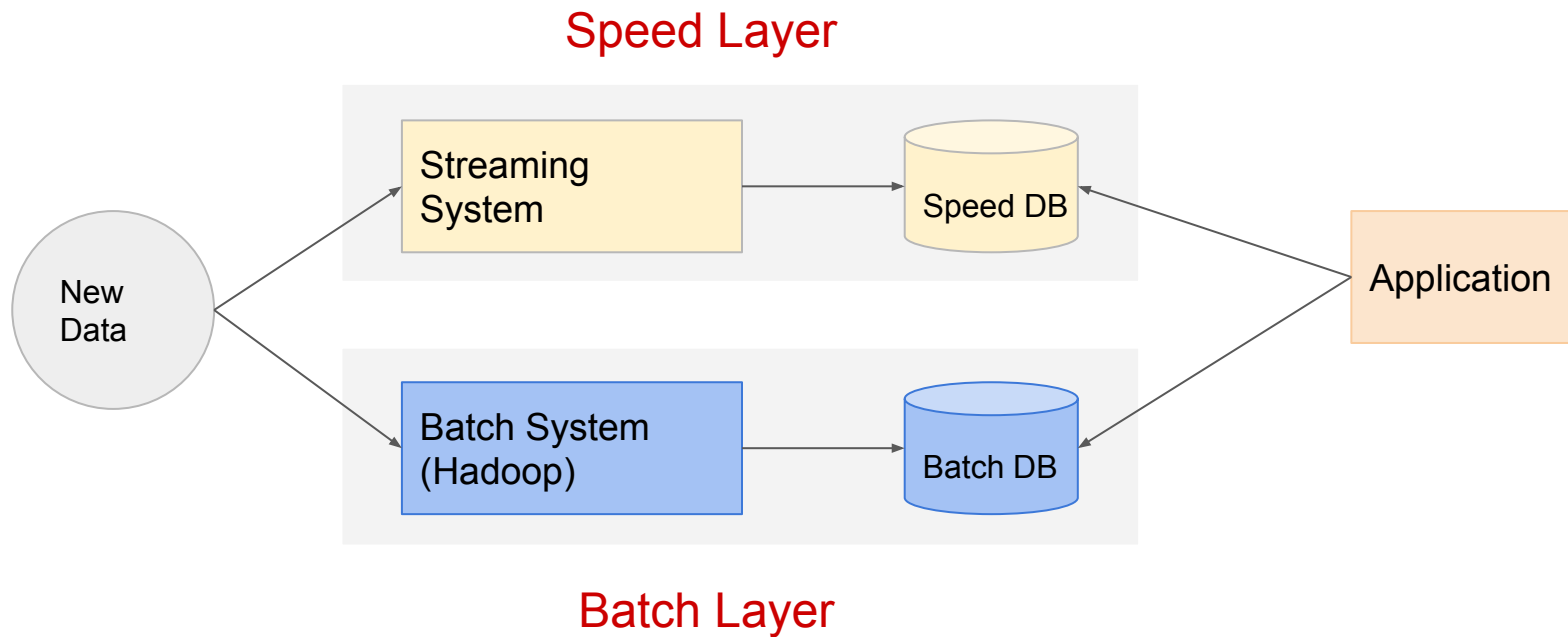
# Contrast with Batch Processing

1. Process Bounded Files - such as files by ingestion time - files by last 15 minutes, last 1 hour, last 1 day
2. Program can go back and forth in data. Do multipass processing.
3. Sessions and joins can span files.
4. Many machine learning algorithms need full batch of data to train.
5. Very high Latency, but very high throughput as well.
  - a. Wait for files to arrive. I.e wait for file window to close.
  - b. Processing the whole file(s) take time.

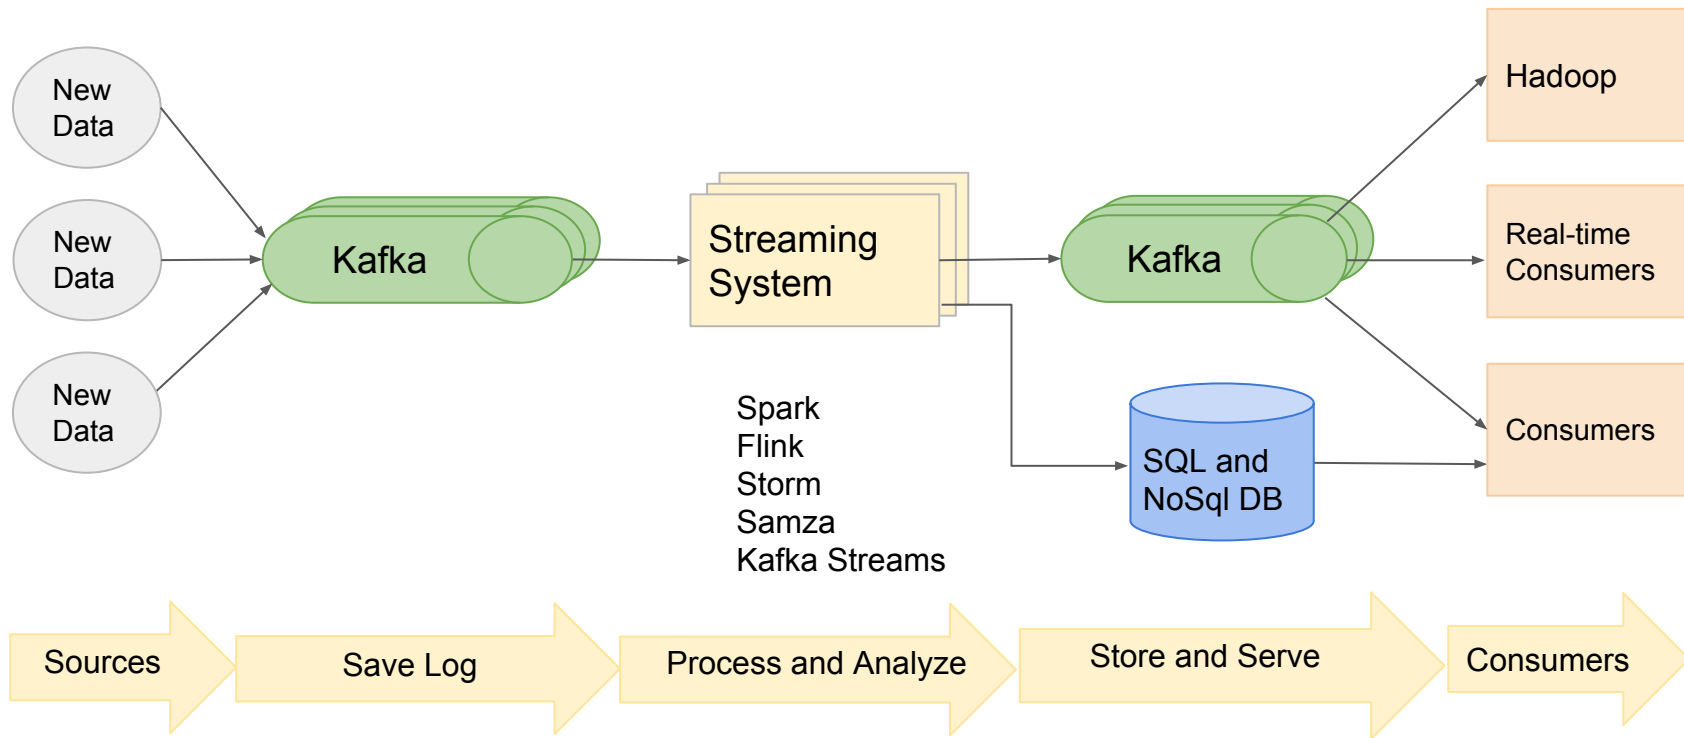
# Streaming Applications

- Joining Clicks and Impressions
- Mobile applications - User activity
- Session based analysis
- Fraud detection
- Industrial IOT
- LinkedIn's Streaming Standardization Platform

# Lambda Architecture



# Streaming Systems Architecture



# What is Streaming Analytics

“ Continuous processing on unbounded data”

“Software that can filter, aggregate, enrich, and analyze a high throughput of data from multiple disparate live data sources and in any data format to identify simple and complex patterns to visualize business in real-time, detect urgent situations, and automate immediate actions.” - Forrester



# Streaming Concepts

## Time

Event Time

Processing Time

## Window

Fixed Window

Sliding Window

Sessions

## Order

Delayed data

Out of order data

## Correctness

Consistency

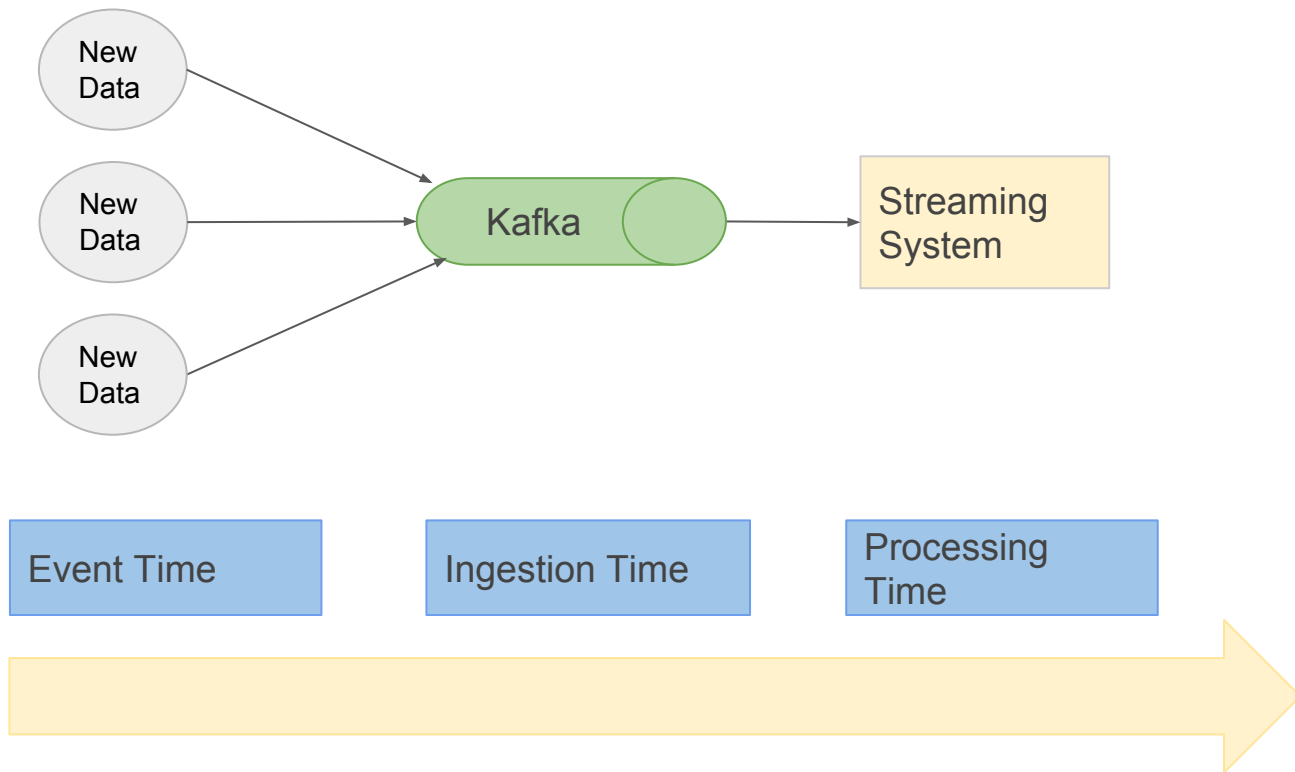
At least Once

Exactly Once


Checkpointing




# Streaming Concepts - Time




# Streaming Concepts - Windows



Fixed Window/  
Tumbling Window








Sliding Window

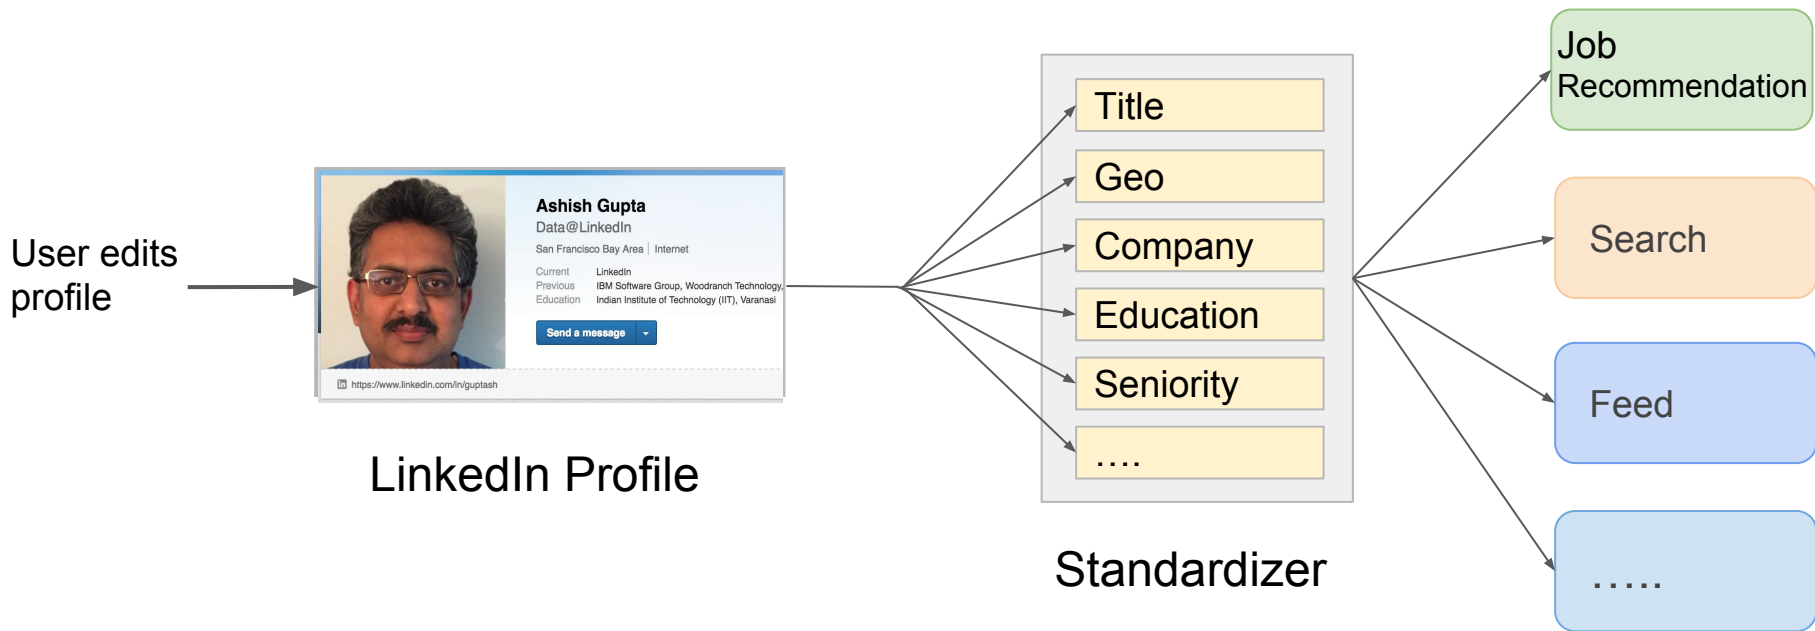


Session Window


# Open Source Streaming Systems

	 <b>SPARK</b>	 <b>Flink</b>	 <b>APACHE STORM</b>	 <b>samza</b>	 <b>Kafka Streams</b>
<b>Processing Model</b>	Mini Batch	Event level	Event level	Event level	Event level
<b>Guarantee</b>	Exactly Once	Exactly Once	At least once	At least once	At least once
<b>State Management</b>	Yes	Yes	No	Yes	Yes
<b>Latency</b>	Medium	Low	Low	Low	Low
<b>Built in primitives</b>	Batch and streaming	Batch and streaming	Low Level API	Low level API	Streaming only
<b>Back Pressure</b>	Yes	Yes	No	via Kafka	via Kafka

# Case Study: LinkedIn Standardization Platform



# Pattern : External Lookup/Stream to Table Join

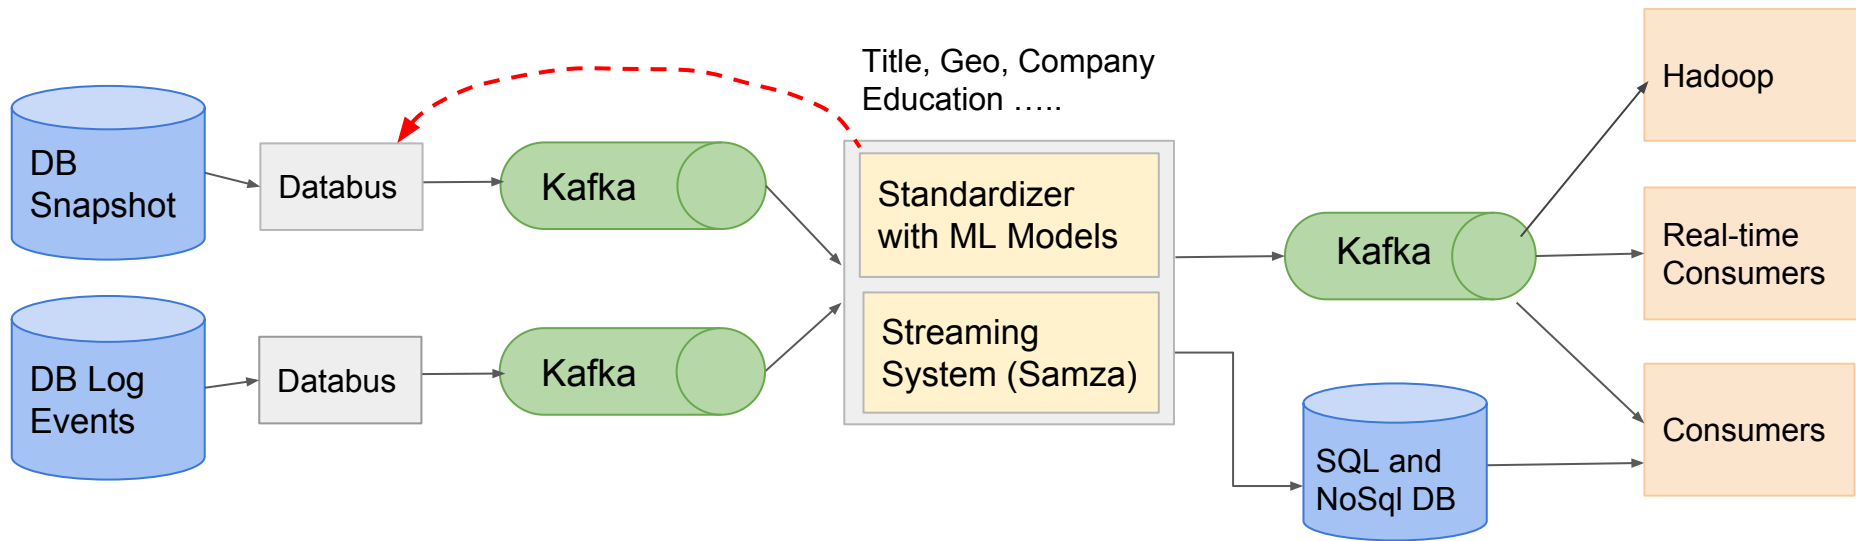


Decide based on size of the data, latency needs and QPS of external systems

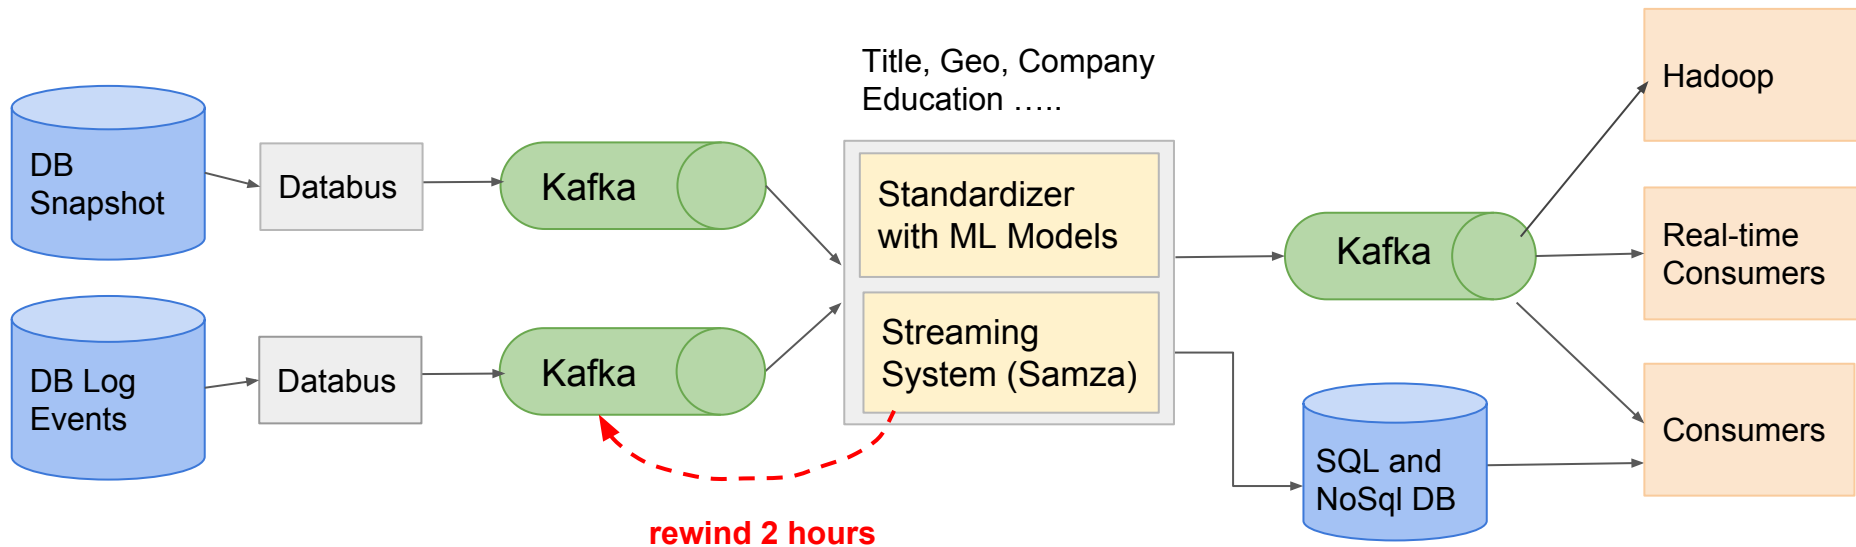
# Pattern: Stream to Stream Joining

- Joins are expensive
- If partitions of two streams not colocated, then expensive shuffle
- Broadcast join if one file is small

# Pattern: Reprocessing



# Pattern: Reprocessing



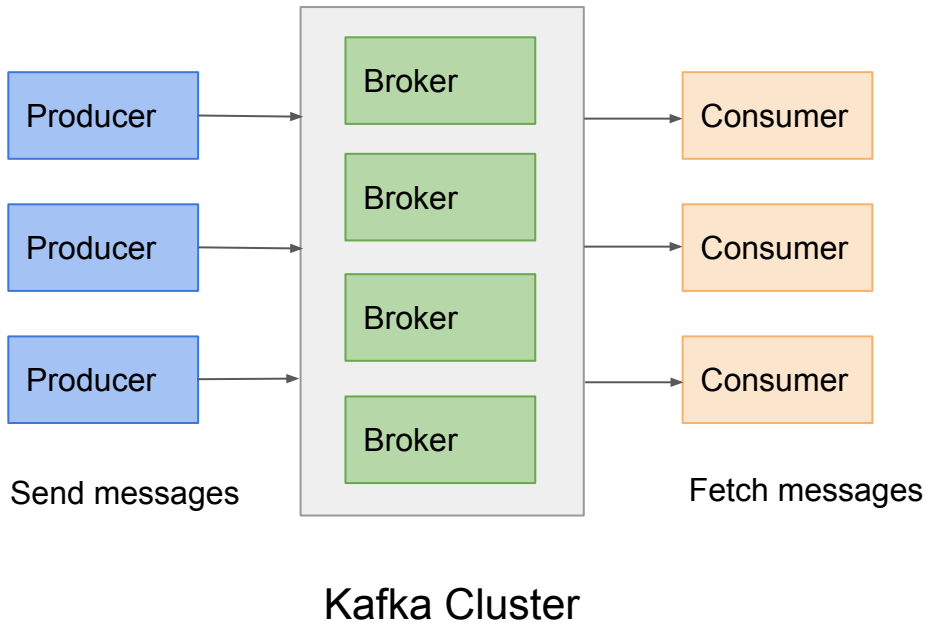


# Apache Kafka

- Highly scalable messaging system
- Distributed commit log
- Developed in LinkedIn back in 2010
- At LinkedIn - more than 1.4 trillion messages per day across over 1400 brokers
- Distributed, partitioned, replicated
- Message retention - based on time and size

# Some Kafka use cases

- Queuing/Messaging
- Metrics
- Auditing
- Logging



# References

- MillWheel: <http://research.google.com/pubs/pub41378.html>
- DataFlow: <http://research.google.com/pubs/pub43864.html>
- Samza: <http://samza.apache.org/>
- Spark Streaming Paper: [Discretized streams](#)
- [Models and Issues in Data Stream Systems](#)

# Contact Us

Ashish Gupta - ahgupta@linkedin.com

<https://www.linkedin.com/in/guptash>

Neera Agarwal - neera8work@gmail.com

<https://www.linkedin.com/in/neera-agarwal-21b9473>